

Mancala

Viktor Kosin und Johanna Beier

04.09.2020

Introduction

Mancala

- ▶ ancient two player game
- ▶ **as vector:** $[6, 6, 6, 6, 6, 6, | 6, 6, 6, 6, 6, 6, | 0, 0]$
- ▶ **Goal:** catch more then half of the beans (37)

Mancala Rules

- ▶ collect all beans of a hole and drop one in each clockwise following hole
- ▶ catch all beans of the last hole, if it contains 6, 4 or 2 beans
- ▶ going backwards: collect beans from all following holes with 6, 4 or 2 beans, if there are no other holes in between
- ▶ game ends if either one player has no more beans or one player catches at least 37 beans
- ▶ total sum of beans: caught beans + beans on own side

MDP

- ▶ Mancala can be represented as a Markov Decision Process (MDP)
- ▶ set of states S , set of actions per state A , action $a \in A$
- ▶ How does the Mancala agent learn to choose the best action?

Reinforcement Learning

- ▶ **Idea:** reward or punish some action
- ▶ **Goal of agent:** maximize total reward
- ▶ **here:** Small reward for catching beans, bigger reward for winning the game
- ▶ use Q-Learning

Q-learning

- ▶ small state space: Q-table
- ▶ replace Q-table by Q-function

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \quad (1)$$

- ▶ agents often need to learn actions that do not lead immediately to a reward
- ▶ allow a small amount of random actions (exploration rate)

bild netz, dass wir verwenden

- ▶ **activation function:** Sigmoidfunction
- ▶ learningrate for the update-weights-function

Backpropagation

- 1. Step** Generate training data: for a given input set an expected output (e.g. with Q-function)
- 2. Step** Calculate for the input $a^{x,1}$:
 - ▶ activation $a^{x,l}$ of layer $l = 2, 3, \dots, L$ by

$$a^{x,l} = \sigma(w^l a^{x,l-1} + b^l)$$

- ▶ Output error $\delta^{x,L}$
- ▶ Backpropagate error to each layer: $\delta^{x,l}$

- 3. Step** Use error of each layer to update weights and biases

play

Training

- ▶ let two agents play against each other and save pairs of actions and rewards
- ▶ update for each boardstate and action the underlying Q-function
- ▶ save each board state and dedicated Q-values as training data
- ▶ feedforward a boardstate to the net
- ▶ $\text{loss} = \text{output} - \text{Q-values}$

Results